

A C-Based Graphics Library for CS1

Eric S. Roberts
Department of Computer Science
Stanford University

NOTE

This paper has been submitted to the Twenty-sixth SIGCSE Technical Symposium on Computer Science Education. Copyright is retained by the author prior to publication release.

ABSTRACT

This paper describes a simple graphics library designed for a CS1 course using ANSI C as its programming language. The library can be implemented easily on a variety of hardware platforms, providing a reasonable level of portability. Implementations currently exist for compilers on the Apple Macintosh, the IBM PC, and Unix workstations; the source code for each of these implementations is publicly available by anonymous FTP from the Roberts.C.CS1 area on host aw.com. In addition, the public distribution includes a fully standard implementation that generates a PostScript representation of the graphical image.

1. INTRODUCTION

Getting college students excited about introductory programming is harder in the 1990s than it was a decade ago. Those who have grown up with personal computers have strong preconceptions about what constitutes an "interesting" application. Modern technology has raised the ante. The simple text-based applications that are the mainstay of most computer science textbooks offer little excitement to the student who is used to modern graphical interfaces. A program that sorts a list of numbers -- as interesting as that problem may be from an algorithmic perspective -- does not impress students who have played computer games for years.

Using a graphics library in the introductory course makes an enormous difference. Even when the homework problems are harder than those assigned in a more traditional course, students approach them with considerably more enthusiasm, which in turn enables them to accomplish more. Captivated by the allure of computer graphics, beginning students have been able to develop extraordinarily sophisticated programs. I have had students complete 1000-line programs in the fourth week of the introductory course.

Beyond fostering excitement, using a graphics library has other pedagogical advantages. In my experience, requiring students to write graphical applications early in the term turns out to be the most effective way to drive home the concepts of parameter passing and stepwise refinement. In constructing graphical figures, students immediately recognize the importance of parameters as they ask themselves, for example, where to put a line or how big to make a circle. Moreover, creating the program structure necessary to display a graphical image encourages students to master the technique of decomposition. Because the graphics library itself provides only a small set of primitive operations such as drawing a straight line, the student quickly learns to assemble lines into rectangles and other

composite figures. As these figures are combined to form still larger structures, the organization of the image suggests a natural decomposition strategy for the program as a whole.

Although these advantages give computer science instructors a strong incentive to integrate graphics into the introductory programming course, doing so is complicated by the following factors:

- o Existing graphics packages are usually designed for a specific implementation platform, which includes not only the underlying hardware but also the operating system, the compiler, and the runtime libraries. An IBM PC program that uses, for example, a graphics package designed for Microsoft Windows will not run on the same machine under DOS, much less on a Macintosh or a Unix system. Because of this high level of system dependence, the instructor who wants to use graphics must choose a particular platform and limit students to the graphical facilities available in that environment.
- o Because there is no single graphical standard that works for all machines, textbooks written to attract the widest possible audience usually avoid the issue by omitting any discussion of graphics. The instructor is therefore forced to develop supplementary course materials that explain how to use the local graphics facilities.
- o Most graphics libraries are designed for use by experts, not introductory students. As a result, those libraries are often too complex for novice programmers to comprehend.

Three years ago, when we started to convert Stanford's introductory course from Pascal to C, we had to face these problems directly. The THINK Pascal system we had been using included a simple graphics library that was well designed for use by beginning students. Unfortunately, the THINK C compiler we chose for the redesigned course offered no corresponding facility. Because our experience had convinced us that using graphics strengthens the introductory course, we decided to design our own C-based graphics library and integrate it with the instructional materials being developed for the course.

The remainder of this paper provides an overview of the graphics library and the strategies used to implement it on a diverse set of programming platforms.

2. INTERFACE DESIGN CRITERIA

The first step in developing the graphics library was to define the interface. In addition to more traditional principles of good interface design, we wanted to develop an interface that met the following criteria:

1. It must be simple. At Stanford, we introduce the graphics library early in the term when students are first learning about functions and procedures. Since the students have relatively few programming tools at their disposal, the graphics interface must not depend on advanced concepts such as pointers, records, or even arrays. Moreover, it should not export so many functions that the novice programmer cannot understand the interface as a whole. Beginning programmers can easily be intimidated by a large interface, even if they

are not required to use all of the functions it contains.

2. It must correspond to student intuition. The conceptual model that underlies the graphics interface must not be so mathematical or so technical that students have trouble understanding its operation. They have usually been exposed to Cartesian coordinate systems in high school mathematics and are now beginning to become familiar with the procedural programming paradigm. The design of the graphics interface should take advantage of that preexisting knowledge, even if the resulting conceptual model is at variance with the underlying model supported by the hardware.
3. It must be powerful enough for students to write programs they think are fun. The principal advantage of the graphics library is that it allows students to undertake programming problems that are exciting enough to capture their imagination. For this to be possible, the interface to that library must export functions powerful enough to generate interesting graphical displays.
4. It must be widely implementable. To ensure that our approach would be relevant to a variety of institutions and not just to Stanford, we believed that it was important to implement the graphics package on several different platforms, particularly those used most often for introductory-level education. We therefore designed the interface so that it did not depend on any specific platform.

In many cases, the individual criteria support one another in that design decisions adopted to satisfy one end up advancing the others. For example, choosing to leave a feature out of the interface results in a library that is not only simpler but more portable, because there are fewer features to implement for each new architecture. In other cases, however, the design criteria trade off against each other, which requires the designer to balance the competing criteria and find an appropriate compromise.

2. EVOLUTION OF THE INTERFACE DESIGN

Fortunately, we have had the opportunity at Stanford to adopt an evolutionary approach to the graphics library design. In each of the four quarters during the year, the Computer Science Department teaches two versions of the introductory programming course: CS106A, for students with little or no prior programming experience, and CS106X, which combines the material in the standard CS1/CS2 curriculum into an intensive one-quarter course for students with more extensive programming backgrounds. Both courses teach ANSI C using a library-based approach [Roberts93] and use the graphics library extensively throughout the course. Thus, we can experiment with a particular library design one quarter and refine it for the next.

We introduced the first version of the graphics library in the fall of 1992, the first quarter in which ANSI C was used for the entire CS106A population. At the time, the graphics library was very simple and included only the following operations:

- o Initialize the library.
- o Move the pen to a specified position on the screen.
- o Draw a line segment with displacements dx and dy.

- o Draw a circle with a given center and radius.

Students were assigned several simple programs using the library and could also take part in an optional graphics contest. The course staff evaluated each entry on the basis of both its artistic merit and its geometrical sophistication. In addition, I invited the students to suggest new features that would have helped them design an even better entry. The following are the most commonly cited suggestions, ranked according to how many times each was cited:

1. Arcs (as opposed to complete circles)
2. Shaded regions
3. Simple animation (by erasing and redrawing a figure)
4. Color
5. Text
6. Mouse input

Having tried to design some graphical displays using the initial version of the library, I was quite sympathetic to the students' concerns. Adding the features they requested would certainly make the interface more powerful and thereby make it more exciting for the better students. Unfortunately, adding those features would also complicate the interface, making it harder for other students to understand. Perhaps more importantly, adding features would make it more difficult to develop implementations for a variety of platforms and therefore compromise the portability criterion.

My solution to this dilemma was to separate the graphics library into two interfaces: a simple `graphics.h` interface providing a minimal set of fundamental tools and a more elaborate `extgraph.h` interface offering the extended capabilities. The textbook developed along with the course [Roberts95] discusses only the simple `graphics.h` interface. Students who seek greater challenges, however, can use the features offered by `extgraph.h`. This design also promotes portability because an implementation that covers only the functions in `graphics.h` is sufficient for all the examples and exercises included in the text. At Stanford, the standard introductory course requires only the `graphics.h` interface, while the accelerated course typically requires the use of `extgraph.h` as well.

The contents of the current versions of each interface are summarized in Figures 1 and 2 at the end of this paper.

4. IMPLEMENTATION STRATEGIES

For the most part, implementing the rendering operations required by the graphics library is not especially difficult. Existing graphics libraries invariably include facilities for drawing lines, displaying text, and filling polygonal regions, which are the only primitives the library requires. Thus, to implement the procedures that actually produce graphical output, all that is needed for each implementation is a small amount of code to translate calls to the `graphics.h` interface into the rendering calls required for the system graphics library.

Similarly, it is also usually a simple matter to create a new window on the screen to use for drawing. Existing graphics libraries usually have high-level tools for window creation, so the `InitGraphics` implementation need only perform the appropriate set of function calls to create and position a new window.

The hard part is implementing the conceptual model used by `graphics.h`, which is fundamentally different from that used in existing graphics libraries. In the typical graphics library supplied with a compiler or operating system, applications are coded using what is generally called the event loop strategy. Each application establishes an initial state and then enters a loop with the following paradigmatic form:

```
void EventLoop(void)
{
    EventT event;

    while (event = WaitForNextEvent()) {
        RespondToEvent(event);
    }
}
```

Calls to `WaitForNextEvent` block until a system event occurs, at which point the operating system delivers the event to the application, which then initiates some appropriate response. These events include, for example, requests to update the window, mouse clicks, and keyboard activity.

While the event-loop paradigm is appropriate for experienced programmers, it makes little sense to students in an introductory course. Those students are used to writing a function called `main` that executes sequentially, one statement at a time. In the first few weeks of a CS1 course, it is unwise to introduce a different paradigm. Thus, students should be able to code graphical applications that start at the beginning of `main` and execute sequentially. For example, to draw a one-inch square centered at the point (1, 1), introductory students expect the program to look like this:

```
main()
{
    InitGraphics();
    MovePen(0.5, 0.5);
    DrawLine(0, 1);
    DrawLine(1, 0);
    DrawLine(0, -1);
    DrawLine(-1, 0);
}
```

This code contains no event loop and has no place for one. Thus, the challenge of implementing the graphics library is to find a way to hide the complexity of the event loop processing inside the library implementation.

Fortunately, this problem is not usually as hard as it appears because most C programming environments already include a solution as part of the standard I/O library implementation. Like the graphics model assumed by the `graphics.h` interface, the I/O model assumed by `stdio.h` is not directly compatible with the event-loop paradigm. For an application that uses `stdio.h` to work, there must be an event loop somewhere that intercepts the keyboard events as they occur. From the perspective of a client of `stdio.h`, however, the event loop is invisible; the program simply calls an input function, such as `getchar` or `scanf`, which eventually returns with the desired input characters.

To preserve the structural simplicity offered by `stdio.h`, microcomputer-based implementations of C typically implement the

standard I/O library so that the required event-loop processing is embedded in the console input routines.[1] Whenever the program waits for input from the standard input device, the library implementation simply enters an event loop, during which it responds to any system events that occur. Keyboard input is queued until a complete line can be returned to the application, but the event loop also processes update requests, so that the program behaves correctly if windows are repositioned on the screen.

The fact that the system must already perform the required event-loop processing makes it easier to implement the graphics library. In most cases, all that is needed is to make the existing event loop used for console I/O also respond to update events for the graphics window. This effect can be achieved in a variety of ways, as illustrated by the following examples:

- o In the Borland C environment for Microsoft Windows, making the graphics window a child of the console window was sufficient to solve the problem, because the parent window automatically delivers events to its children.
- o In the THINK C environment on the Macintosh, it was necessary to make a dynamic patch to the console event loop so that it also called the appropriate update procedures in the graphics library.
- o The X Windows implementation for Unix required a different approach. In this implementation, the function `InitGraphics` calls `fork` to create two parallel threads, connected by a communication pipe. One of those threads returns to execute the application code. The other enters an event loop waiting either for X events that require a response or for commands from the application program sent over the communication pipe.

In addition to these platform-specific implementations, the public distribution site on `aw.com` also includes a fully standard version that does not actually display a graphical image but instead writes a PostScript file suitable for printing. The existence of this implementation ensures that it is always possible to use the simple `graphics.h` interface even on platforms for which no specialized implementation exists.

5. CONCLUSIONS

Our experience over the last two years at Stanford suggests that using a graphics library in an introductory course enhances student interest and helps reinforce several essential programming concepts. We have also demonstrated that it is possible to design a graphics library that can be implemented in a relatively portable way.

REFERENCES

[Hilburn93] Thomas B. Hilburn, "A top-down approach to teaching an introductory computer science course," SIGCSE Bulletin, March 1993.

[House94] Donald House and David Levine, "The art and science of computer graphics: a very depth-first approach to the non-majors course," SIGCSE Bulletin, March 1994.

[Papert80] Seymour Papert, *Mindstorms*, New York: Basic Books, 1980.

[Roberge92] James Roberge, "Creating programming projects with visual impact," SIGCSE Bulletin, March 1992.

[Roberts93] Eric S. Roberts, "Using C in CS1: Evaluating the Stanford experience," SIGCSE Bulletin, March 1993.

[Roberts95] Eric S. Roberts, The Art and Science of C: A Library-Based Approach, Reading, MA: Addison-Wesley, 1995.

FOOTNOTES

[1] In THINK C, for example, this operation is performed as part of the implementation of the console I/O package console.c. In the Borland C/C++ system, the same functionality is provided by the EasyWin package.

Figure 1. Functions in graphics.h

<code>InitGraphics();</code>	This procedure initializes the package and creates the graphics window on the screen. The call to the <code>InitGraphics</code> function is typically the first statement in <code>main</code> .
<code>MovePen(x, y);</code>	This procedure moves the current point to the position (x, y) without drawing a line.
<code>DrawLine(dx, dy);</code>	This procedure draws a line extending from the current point by moving the pen dx and dy inches along the respective coordinate axes.
<code>DrawArc(r, start, sweep);</code>	This procedure draws a circular arc, which always begins at the current point. The arc itself has radius r, and starts at the angle specified by the parameter start, relative to the center of the circle. This angle is measured in degrees counterclockwise from the 3 o'clock position along the x-axis, as in traditional mathematics. The fraction of the circle drawn is specified by the parameter sweep, which is also measured in degrees. If the value of sweep is positive, the arc is drawn counterclockwise from the current point; if sweep is negative, the arc is drawn clockwise. The current point at the end of the <code>DrawArc</code> operation is the final position of the pen along the arc.
<code>width = GetWindowWidth();</code> <code>height = GetWindowHeight();</code>	These functions return the width and height, in inches, of the drawing area of the graphics window.
<code>cx = GetCurrentX();</code> <code>cy = GetCurrentY();</code>	These functions return the current x- and y-coordinates of the pen. Like all coordinates in the graphics package, these values are measured in inches relative to the origin.

Figure 2. Selected functions in extgraph.h

<code>DrawEllipticalArc (rx, ry, start, sweep);</code>	This procedure is similar to <code>DrawArc</code> in the basic graphics library, but draws an elliptical rather than a circular arc.
<code>StartFilledRegion(density); EndFilledRegion();</code>	These procedures are used to draw a filled region. The boundaries of the region are defined by a set of calls to the line and arc drawing functions that represent the boundary path. The parameter <code>density</code> is a floating-point value that represents the pixel density to be used for the fill color.
<code>DrawTextString(s);</code>	This procedure draws the text string <code>s</code> at the current pen position. The pen position is moved so that it follows the final character in the string. The interface also contains functions to set the font, size, and style of the text.
<code>SetPenColor(color); DefineColor(name, r, g, b,);</code>	The <code>SetPenColor</code> function sets the graphics package to draw in some color other than the default black. Colors are identified by name, which makes it easier to save and restore the current color. By default, only a small set of standard colors are defined, but <code>DefineColor</code> allows clients to define new ones by specifying their RGB values.
<code>x = GetMouseX(); y = GetMouseY(); flag = MouseButtonIsDown();</code>	These functions make it possible for the client to use the mouse to design graphical user interfaces. Given these primitives, it is easy to design interfaces that provide portable implementations of buttons, menus, and other standard interactors.
<code>Pause(seconds);</code>	The <code>Pause</code> function forces the system to redraw the screen and then waits for the specified number of seconds. This function can be used to animate the displays.